



# Operating Systems


## Scheduling Algorithms

Seyyed Ahmad Javadi

[sajavadi@aut.ac.ir](mailto:sajavadi@aut.ac.ir)

Fall 2022

# Classes



	Saturday	Monday	Saturday	Monday
	Green	Green	Red	Red
	Red	Red	Red	Red
	Green	Green	Green	Green
	Green	Red	Light Green	Light Green
	Light Green	Light Green	Light Green	Light Green
	Light Green	Light Green	Light Green	Light Green
	Light Green	Light Green	Light Green	Light Green
	Light Green	Light Green	Light Green	Light Green

---

# First-Come, First-Served



# First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$
- The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$P_2, P_3, P_1$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- **Much better than previous case**

# FCFS Scheduling and Convoy effect

---

- **Short process behind long process.**
  - Consider one CPU-bound and many I/O-bound processes.



- **What is the important side-effect?**

# FCFS Scheduling and Convoy Effect (Cont.)

- Short process behind long process.
  - Consider one CPU-bound and many I/O-bound processes.



- What is the side-effect?
  - Results in **lower CPU and device utilization** than might be possible if the shorter **processes were allowed to go first.**

---

# Shortest-Job-First






# Shortest-Job-First (SJF) Scheduling

---

- **Associate** with each process the length of its **next CPU burst**.
  - Use these lengths to schedule the process with the shortest time.



Ready Queue

# Shortest-Job-First (SJF) Scheduling (cont.)

---


- **SJF is optimal**
  - Gives minimum average waiting time for a given set of processes.
  
- Preemptive version called **shortest-remaining-time-first**
  
- **The difficulty is knowing the length of the next CPU request**
  - Could ask the user
  - Estimate (we do not cover this in the class and the exams)



# Example of SJF

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

## ■ SJF scheduling chart



## ■ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

---

# Round-Robin



# Round Robin (RR)

---

- Each process gets a small unit of CPU time (**time quantum  $q$** )
  - Usually 10-100 milliseconds.
  - After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ :
  - Each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once.
  - No process waits more than  $(n-1)q$  time units.



# Round Robin (RR) (cont.)

---


- Timer *interrupts* every quantum to schedule next process
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high



# Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better **response**
- $q$  should be large compared to context switch time
  - $q$  usually 10 milliseconds to 100 milliseconds,
  - Context switch < 10 microseconds

# Time Quantum and Context Switch Time





# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

A rule of thumb is that  
80% of CPU bursts should  
be shorter than  $q$

---

# Priority Scheduling



# Priority Scheduling

---

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority  
(**smallest integer  $\equiv$  highest priority**)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling where priority is the ***next predicted CPU burst time***



# Priority Scheduling

---

## ■ Problem ≡ Starvation

- Low priority processes may never execute


## ■ Solution ≡ Aging

- As time progresses increase the priority of the process

# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

## ■ Priority scheduling Gantt Chart




## ■ Average waiting time = 8.2

# Priority Scheduling w/ Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	4	3
$P_2$	5	2
$P_3$	8	2
$P_4$	7	1
$P_5$	3	3

- Run the process with the highest priority. Processes with the same priority run round-robin
- Gantt Chart with time quantum = 2




---

# Multilevel Queue Scheduling



# Multilevel Queue

- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue!






# Multilevel Queue

- Prioritization based upon process type

highest priority



lowest priority

# Multilevel Feedback Queue


---

- A process can move between the various queues.
- Multilevel-feedback-queue defined by the following parameters:
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to determine when to upgrade a process
  - Method used to determine when to demote a process
  - Method used to determine which queue a process will enter when that process needs service
- **Aging** can be implemented using multilevel feedback queue



# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS



# Example of Multilevel Feedback Queue (cont.)

## ■ Scheduling

- A new process enters queue  $Q_0$  which is served in RR
  - ▶ When it gains CPU, the process receives 8 milliseconds
  - ▶ If it does not finish in 8 milliseconds, the process is moved to queue  $Q_1$
- At  $Q_1$  job is again served in RR and receives 16 additional milliseconds
  - ▶ If it still does not complete, it is preempted and moved to queue  $Q_2$

