



Operating Systems

Threads and Concurrency

Seyyed Ahmad Javadi

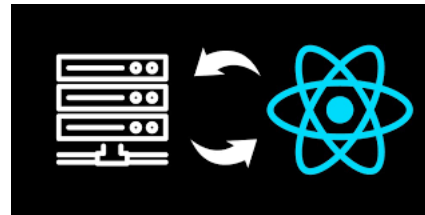
sajavadi@aut.ac.ir

Fall 2022

Motivation

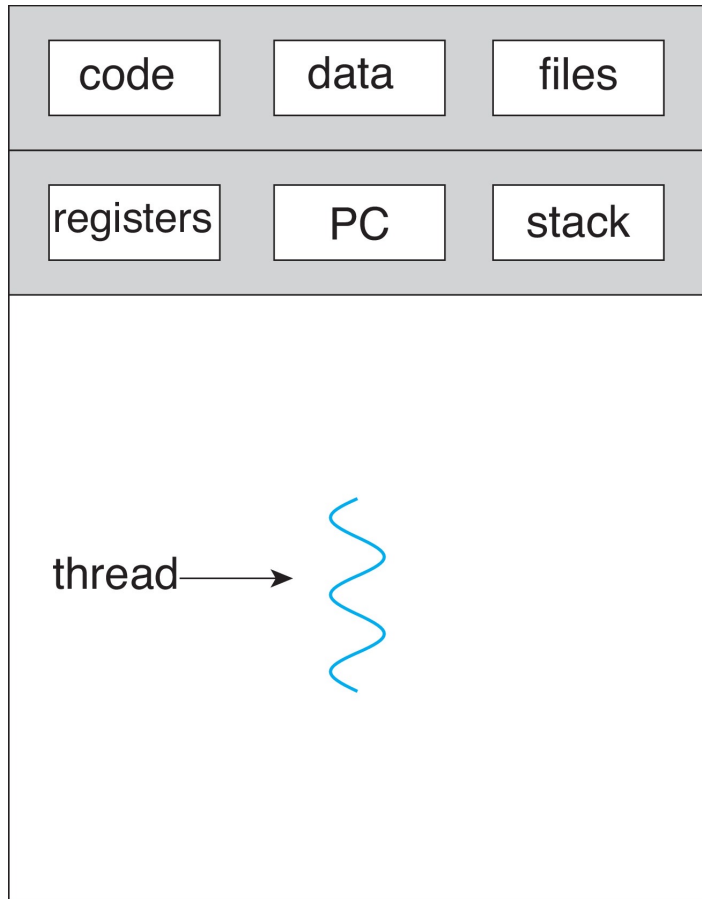
- Most modern applications are multithreaded
- Multiple tasks with the application can be implemented by threads

- Update display
- Fetch data
- Spell checking

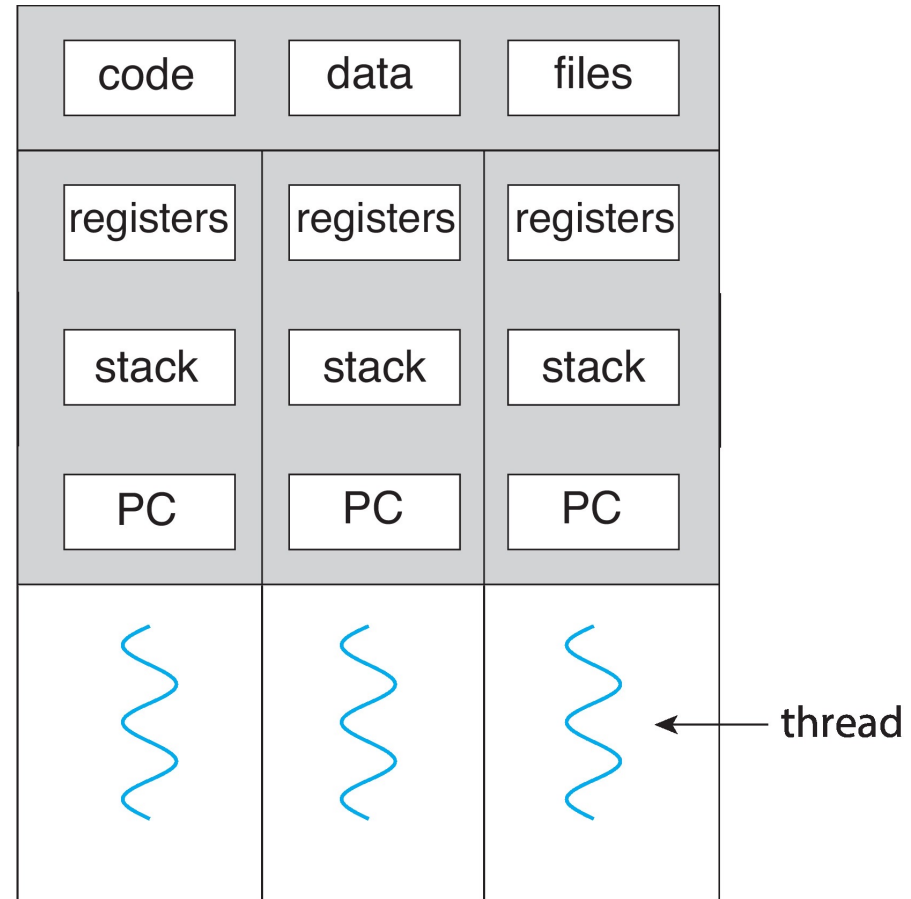


- **Process** creation is **heavy-weight** while **thread** creation is **light-weight**
- **Kernels are generally multithreaded**

Single and Multithreaded Processes

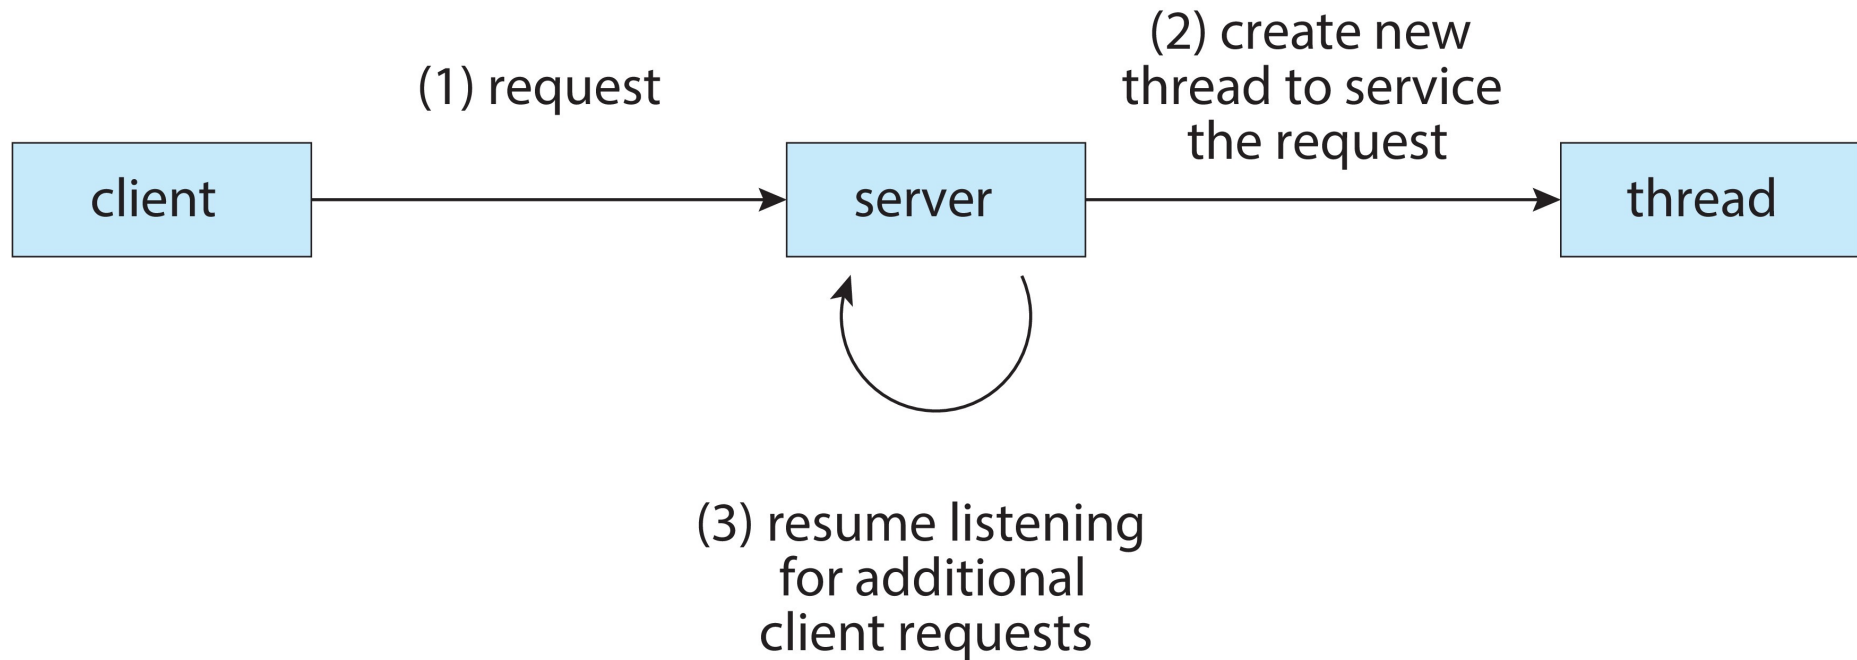


single-threaded process



multithreaded process

Multithreaded Server Architecture



Benefits

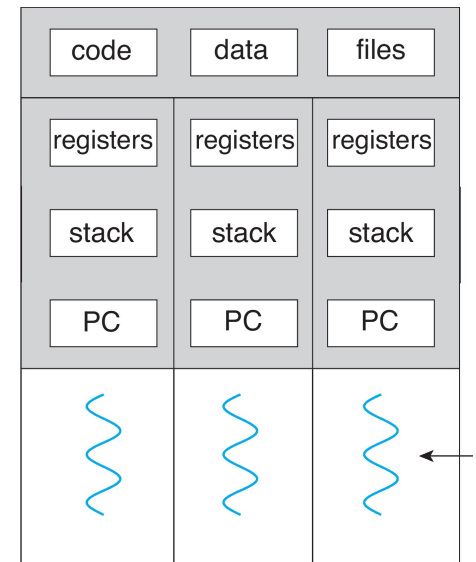
■ Responsiveness

- Allow continued execution if part of process is blocked
- Especially important for user interfaces



■ Resource Sharing

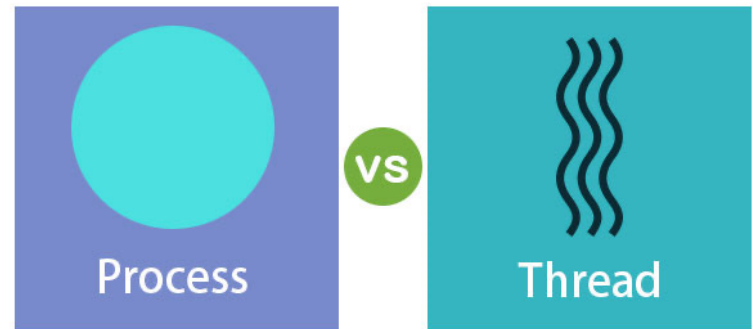
- Threads **share resources** of process
- **Easier** than shared memory or message passing.



Benefits (cont.)

■ Economy

- Cheaper than process creation
- Thread switching lower overhead than context switching.



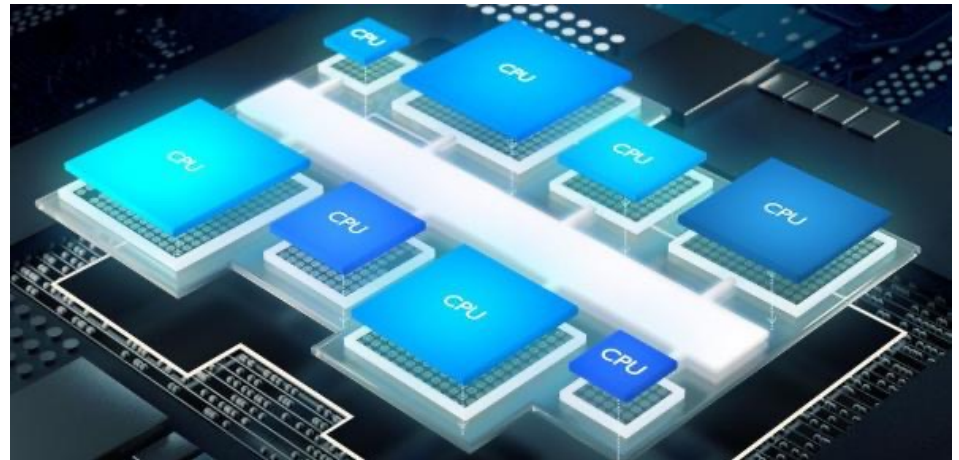
■ Scalability

- Process can take advantage of multicore architectures.



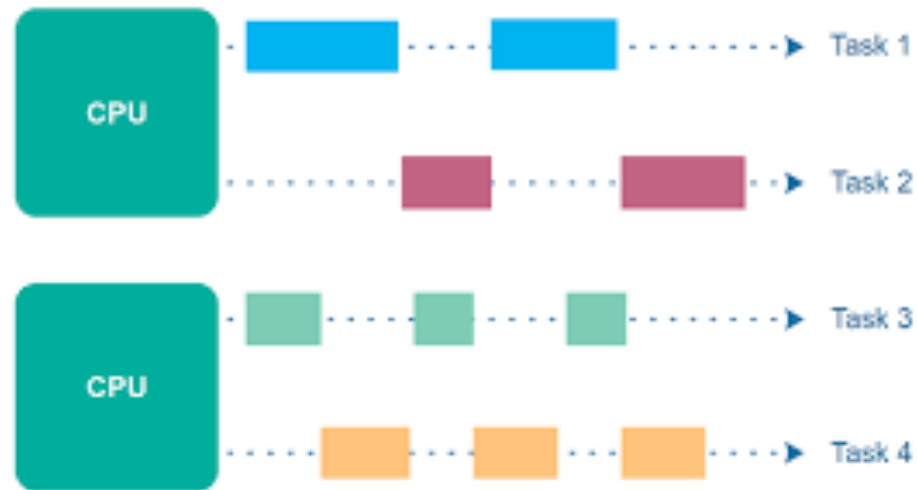
Multicore Programming

- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance: ensuring that the tasks perform equal work of equal value.
 - Data splitting
 - Data dependency
 - Testing and debugging



Multicore Programming (cont.)

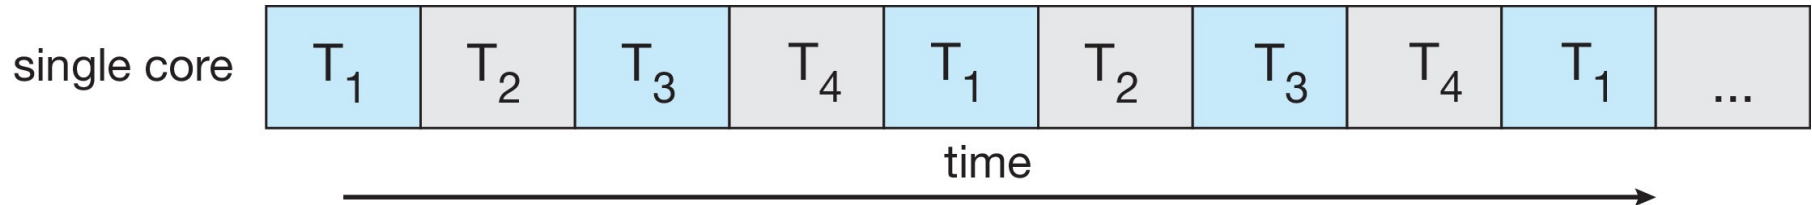
- **Parallelism** implies performing more than one task **simultaneously**



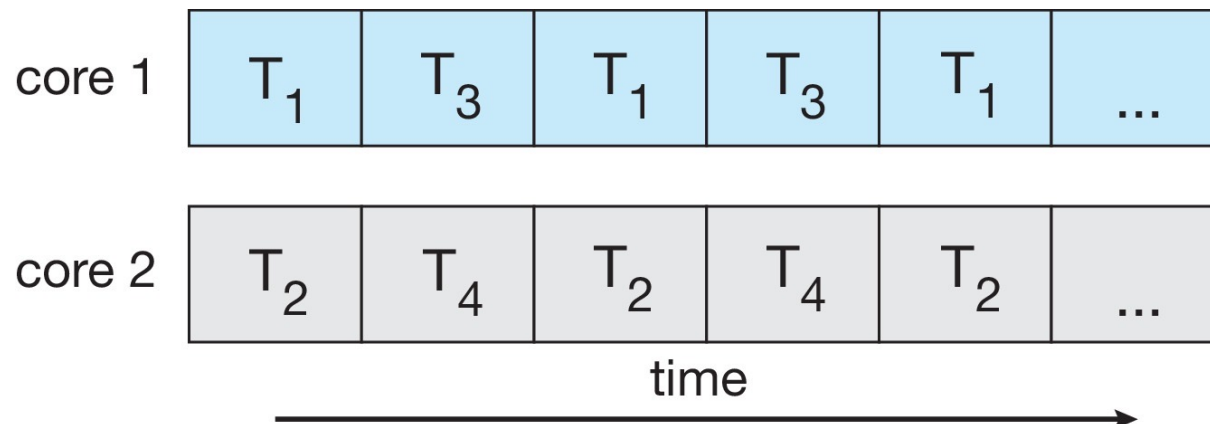
- **Concurrency** supports more than one task **making progress**
 - Single processor or core, scheduler providing concurrency

Concurrency vs. Parallelism

- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**



Multicore Programming-Types of Parallelism

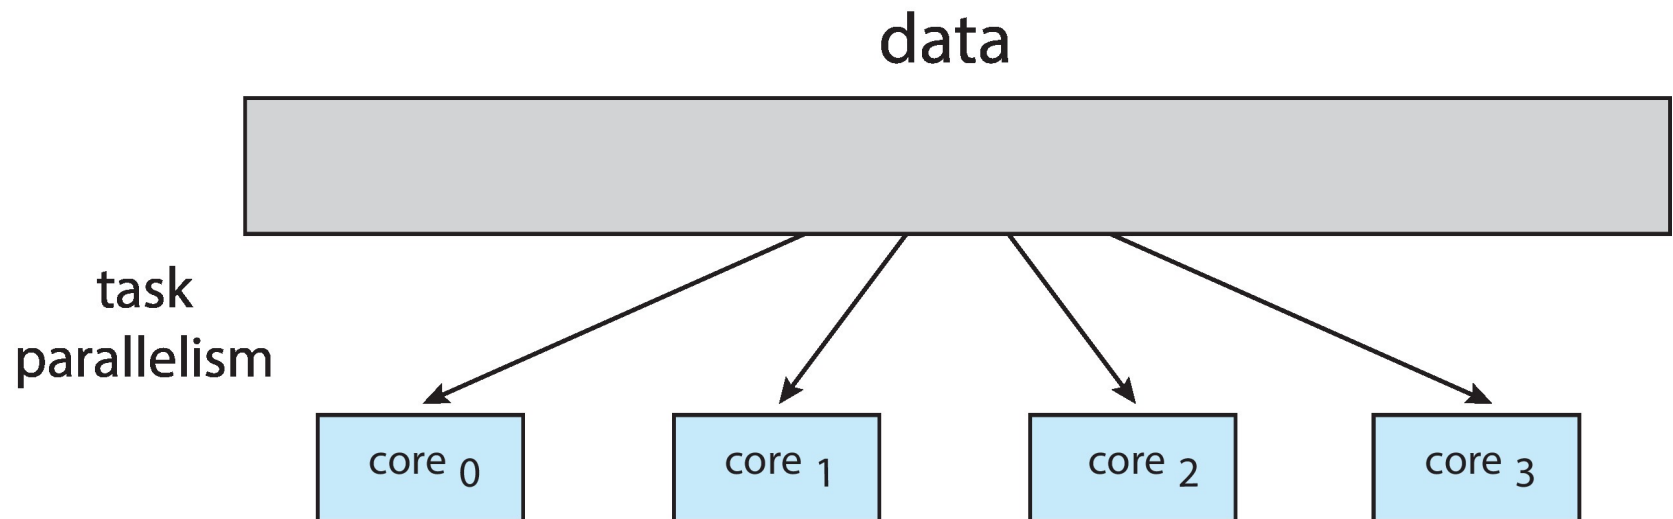
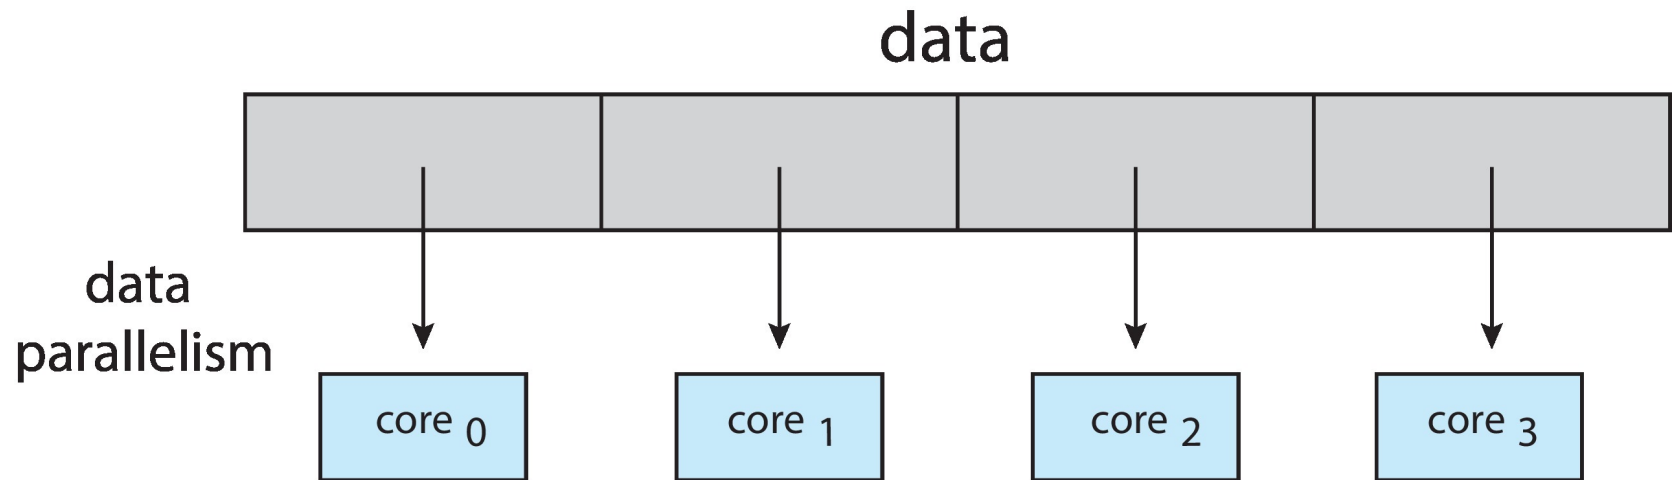
■ Data parallelism

- Distributes subsets of the same data across multiple cores, same operation on each
- Example: summing the contents of an array of size N.

■ Task parallelism

- Distributing threads across cores, each thread performing unique operation
- Example: Unique statistical operation on the array of elements.

Data and Task Parallelism

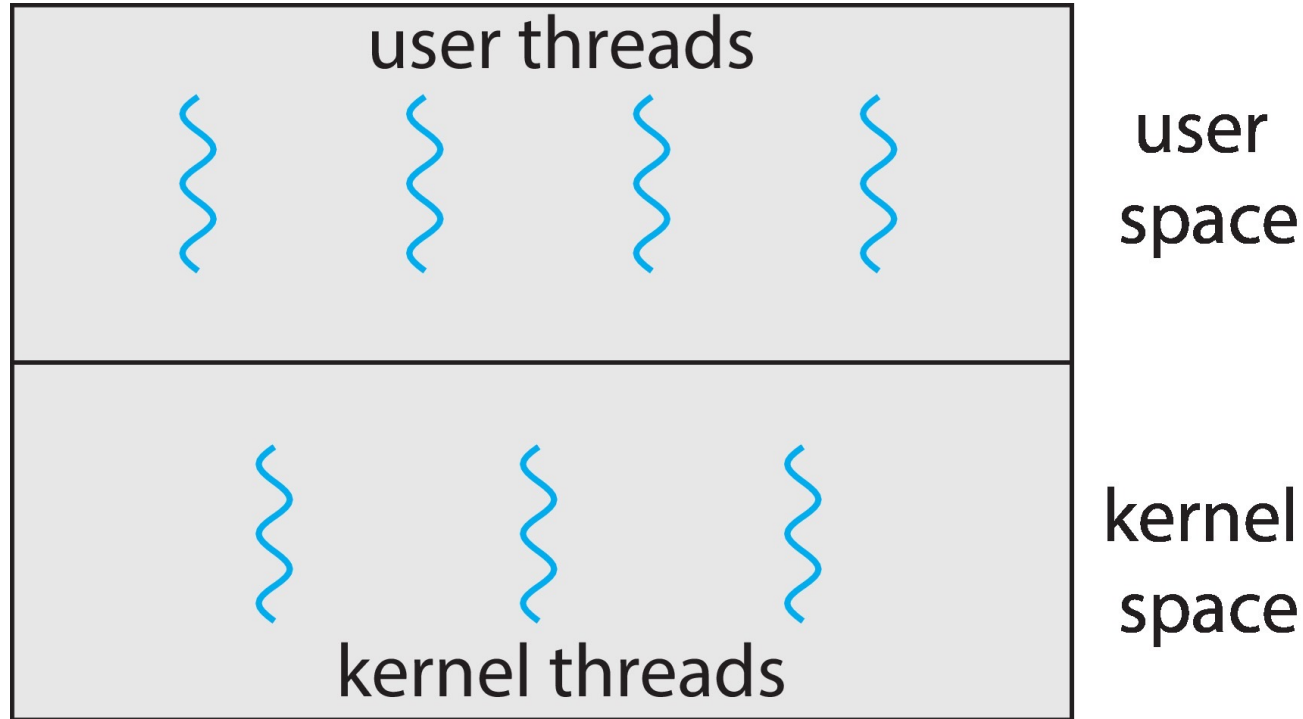


User Threads and Kernel Threads

- **User threads:** management done by user-level threads library.
- Three primary thread libraries:
 - POSIX Pthreads
 - Windows threads
 - Java threads
- **Kernel threads:** supported by the Kernel
 - Examples – virtually all general -purpose operating systems, including: Windows, Linux, Mac OS X, iOS, Android



User and Kernel Threads



Additional review: <https://www.geeksforgeeks.org/difference-between-user-level-thread-and-kernel-level-thread/>

Multithreading Models

- How to map user threads to kernel threads?
- Many-to-One
- One-to-One
- Many-to-Many

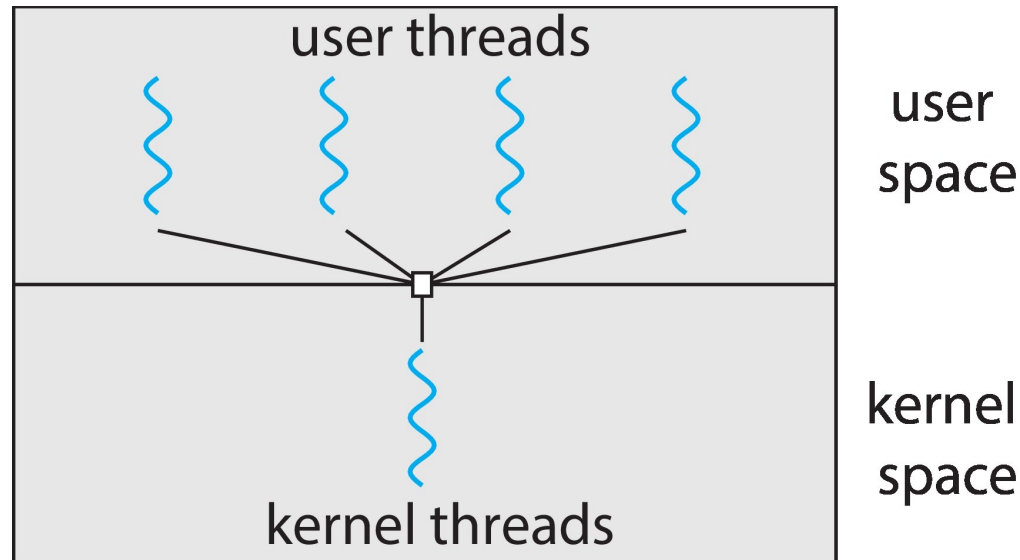
Additional review:

<https://stackoverflow.com/questions/14791278/threads-why-must-all-user-threads-be-mapped-to-a-kernel-thread>




Many-to-One

- Many user-level threads mapped to single kernel thread
- Thread management is done by the thread library in user space
 - So it is efficient



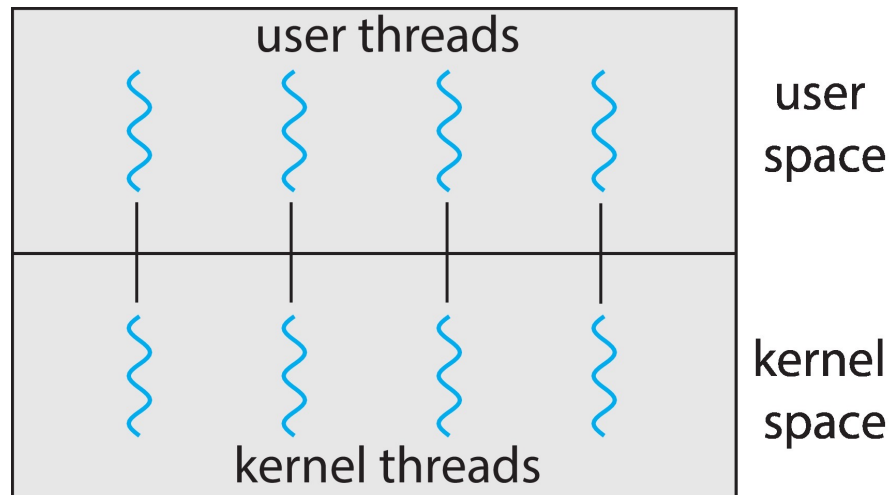
Many-to-One (cont.)

- One thread blocking causes all to block
- ***Multiple threads may not run in parallel on multicore system***
 - Because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads



One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread




One-to-One (cont.)

- **More concurrency than many-to-one**
- Number of threads per process may be restricted due to overhead
- Examples

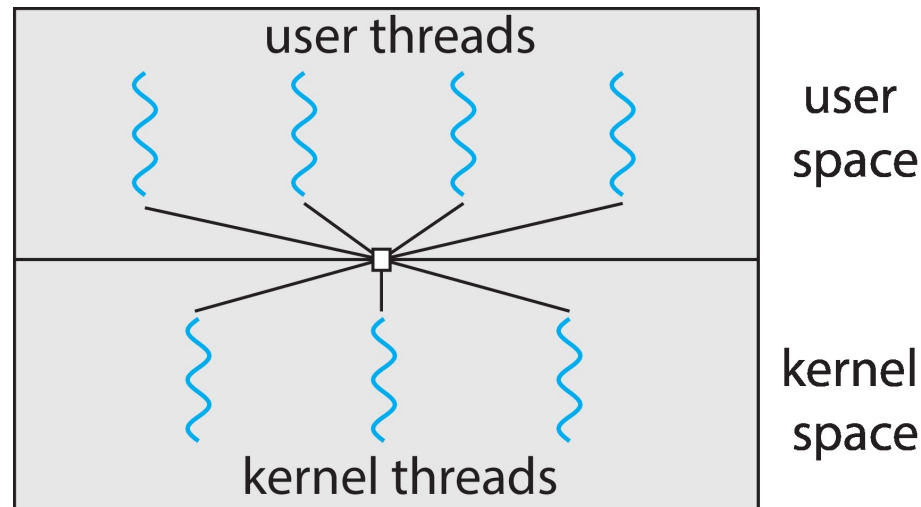
- Windows

- Linux



Many-to-Many Model

- Many user level threads to be mapped to many kernel threads
- Operating system can create a sufficient number of kernel threads



- Examples
 - Windows with the *ThreadFiber* package
 - Otherwise not very common

Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread

